

Anthony Poerio  
[adp59@pitt.edu](mailto:adp59@pitt.edu)  
University of Pittsburgh  
CS1571 – Artificial Intelligence  
Homework #03  
Flappy Bird AI – Report

## Project Overview

For our third assignment in Dr. Hwa's AI class, our goal was to create an intelligent agent able to automatically play **Flappy Bird**. The agent we were tasked with creating should learn how to play Flappy Bird via a **Reinforcement Learning** technique called **Q-Learning**.

The version of **Flappy Bird** that we were tasked with updating is a Python implementation created by Timo Wilken and available for download directly at: <https://github.com/TimoWilken/flappy-bird-pygame>. This Flappy Bird version is implemented using the **PyGame** library.

More specifically, **Flappy Bird** is simple computer game in which the user controls a small bird, and at any given time is given two choices: Stay, or Jump. The user must select the best option at any given time in order to pass obstacles in the bird's path, thereby scoring points.

Our AI is tasked with playing Flappy Bird automatically. This write-up documents the decisions I made in programming my Flappy AI.

## Framework Approach

The first challenge I had in implementing the framework for my Flappy AI was determining exactly how the game worked, in its original state.

By using the debugger and stepping through the game's code during some trial runs, I was able to figure out where key decisions were made, how data flowed into the game, and exactly where I would need to position my agent.

At its basic level, I created an "Agent" class, and passed that class into the running game code. Then, at each loop of the game, I examined the variables available to me, and then passed a 'MOUSEBUTTONUP' command to the PyGame event queue whenever the AI decided to jump. Otherwise, I did nothing.

## State Representations

From there, the next step was determining a way to model the problem. I decided to use follow the basic guidelines [outlined by Sarvagya Vaish, here](#).

First, I discretized the space in which the bird sat, relative to the next pipe. I was

able to get pipe data by accessing the **pipe object** in the original game code. Similarly, I was able to get bird data by accessing the **bird object**.

From there, I could determine the location of the **bird** and the **pipes** relative to each other. I discretized this space as a 16x16 grid, with the following parameters:

```
# first value in state tuple
height_category = 0
dist_to_pipe_bottom = pipe_bottom - bird.y
if dist_to_pipe_bottom < 8: # very close
    height_category = 0
elif dist_to_pipe_bottom < 20: # close
    height_category = 1
elif dist_to_pipe_bottom < 125: # mid
    height_category = 2
elif dist_to_pipe_bottom < 250: # far
    height_category = 3
else:
    height_category = 4

# second value in state tuple
dist_category = 0
dist_to_pipe_horz = pp.x - bird.x
if dist_to_pipe_horz < 8: # very close
    dist_category = 0
elif dist_to_pipe_horz < 20: # close
    dist_category = 1
elif dist_to_pipe_horz < 125: # mid
    dist_category = 2
elif dist_to_pipe_horz < 250: # far
    dist_category = 3
else:
    dist_category = 4
```

Using this methodology, I created a **state tuple** that looked like this:  
(**height\_category**={0,1,2,3,4}, **dist\_category**={0,1,2,3,4} , **collision**=True/False)

Then, each iteration of the game loop, I was able to determine the bird's relative position, and whether it had made a collision with the pipes or not.

If there was no collision, I issued a reward of +1.

If there was a collision, I issued a reward of -1000.

I tried many different state representations here, but mostly it was matter of determining an optimal number of grid spaces and the right parameters for those spaces.

Initially, I started with a 9x9 grid, but moved to 16x16 because I got to a point in 9x9 where I just couldn't make any more learning progress.

Very generally, we want to have a **tighter grid around the pipes**, as this is where most collisions happen. And we want a **looser grid as we move outwards**. This

seemed to give me the best results, as we need different strategies at different locations on the grid.

## Exploration Approaches

My next task was implementing an exploration approach.

Because we have only two choices at any given state (JUMP—or—STAY), implementing exploration was relatively simple.

I started out with a high **exploration factor** (I used  $1/\text{time\_value}+1$ ), and then I generated a random number between [0,1). If the random number was less than the exploration factor, then I explored.

Over time the exploration factor got lower, and therefore the AI explored less frequently.

Exploration essentially consisted of flipping a fair coin (generating a Boolean value randomly).

- If true → then I chose to JUMP.
- If false → I chose to STAY.

The main problem I encountered with this method is that the exploration factor was very at the beginning, and sometimes choices were made that were not representative of actual situations that the bird would encounter in 'true' gameplay.

BUT, because these decisions were made earlier, they were weighted more heavily in the overall Q-Learning algorithm.

This isn't ideal, but exploration is necessary, and overall the algorithm works well. So it wasn't a large problem, overall.

## Learning Rates & Their Impact

The first learning rate I tried was  **$\alpha = (1/\text{time}+1)$** . However, this gave very poor results in practice.

This is because time is NOT the most important factor in determining a strategy from any given state. Rather, it is how many times we've been to that state.

The problem is that we make extremely poor choices at the beginning of the game (because we simply don't know any better). But with  **$\alpha = (1/\text{time}+1)$** , the results of these these poor choices are weighted the most highly.

Once I changed the learning factor to  **$\alpha = 1/N(s,a)$** , I immediately saw dramatically better results. (That is, where  $N(s,a)$  tracks how many times we've been in a given state and performed the same action.)

# Training

My final, “Smart” bird is the result of about **4 hours of training**.

I don’t actually think there would be a way to make the training more efficient, aside from speeding up the gameplay in some way.

Overall, I the results I received from the investment of time I put it in reasonable.

Given more time, I would probably ***discretize the space even more finely*** (maybe a 25x25 or 36x36 grid) – so that I could find even more optimal strategies from a more fine-tuned set of positions in the game-space.

# How to Use my Smart Bird

To use my smart bird, simply take the following steps:

1. cd into a directory containing my source code
2. Ensure that this directory includes the file named ‘**qdata.txt**’
3. Run the command:
  - a. **python flappybird.py “qdata.txt”**
4. Watch Flappy crush it. (the game will run 10x)

# Citations

I consulted the following resources to implement my AI:

- **Implementation ideas** based on discussion here:  
<http://sarvagyaish.github.io/FlappyBirdRL/>
- **Code** created from Dr. Hwa’s Pseudo-Code here:  
<https://docs.google.com/document/d/1r1X8lvDLYkk1ztHNbcZsd1Lfp2gqP4AZvivKI509G7Q/edit>
- **For understanding what I was doing better:**  
<http://mnemstudio.org/path-finding-q-learning-tutorial.htm>